

μ Caspian

Technical Documentation
Parker Mitchell

Caspian Platform

- A Platform for developing Spiking Neural Networks which may be deployed onto neuromorphic architectures
- Includes:
 - Fixed-increment discrete event simulator (C++)
 - Network creation and manipulation API (C++ / Python)
 - Unified simulation & hardware API (C++ / Python)
- Currently building out novel neuromorphic hardware architectures
- μCaspian is the first of those architectures

μ Caspian v1

- 256 neurons, 4096 synapses (on a Lattice iCE40 up5k FPGA)
- Integrate and Fire Neuron model
 - 8-bit thresholds, 16-bit charge state, 4-bit axonal delay (0-15 cycles)
- Sparse directed graph connectivity embedding
 - Outgoing synapses are stored as a contiguous range for each axon/neuron
 - Each synapse stores the id of its target neuron as well as an 8-bit signed weight value
- Deterministic, event-driven execution model
 - Exact match with software simulation
 - Each time step requires a variable execution time dependent on network activity
- Integrated with software ecosystem
 - Caspian simulator, EONS training algorithm, TENNLab framework, etc.
 - Hardware is usable through a C++ or Python API

The µCaspian Architecture

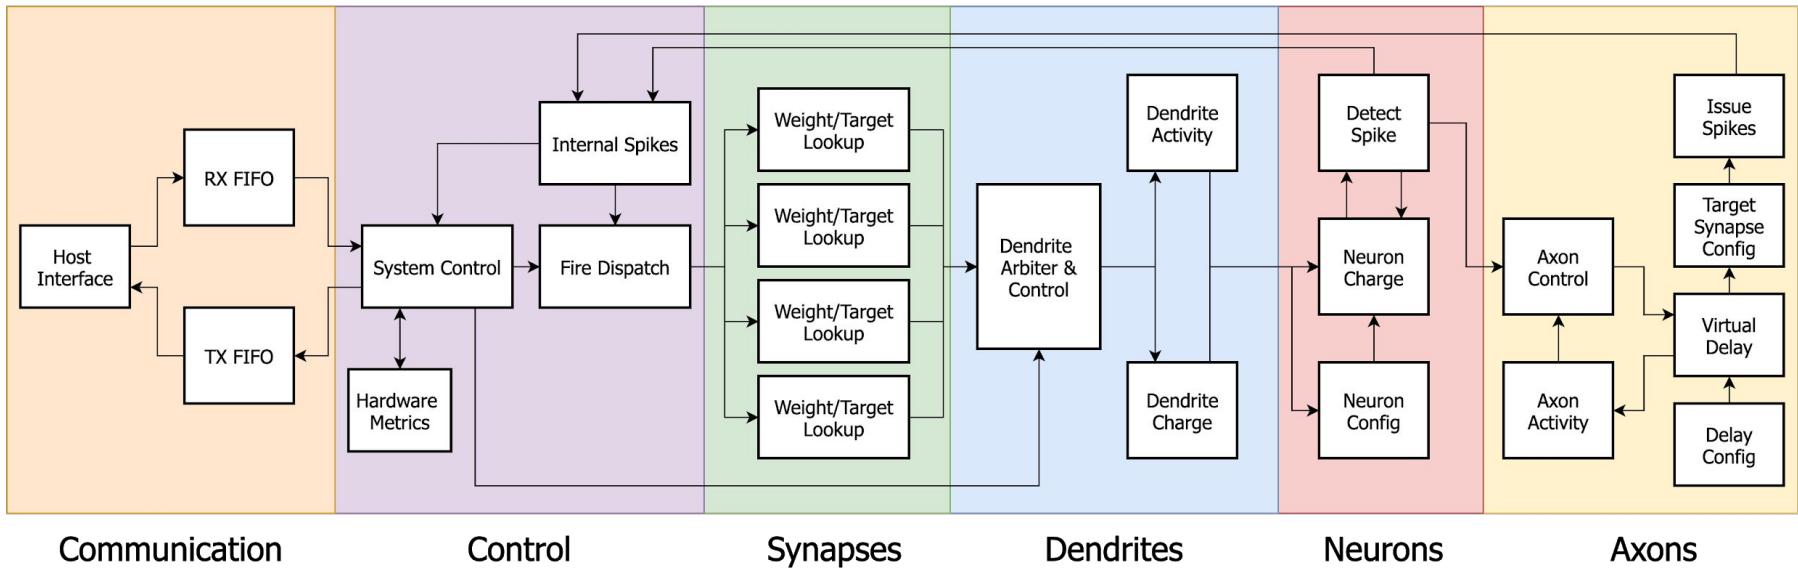
Architecture Concepts

“Pipeline of pipelines” - The architecture uses a dataflow pipeline of modules in which each module may be internally pipelined. Inter-module communication uses a standard 2-bit handshake flow control protocol (ready/valid).

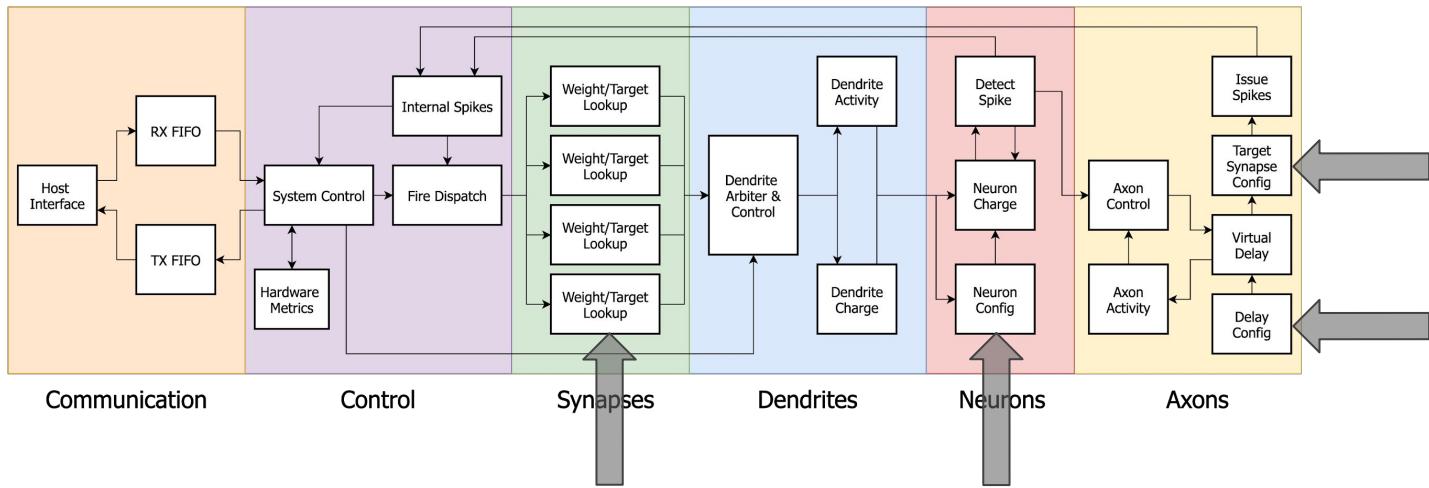
Event-driven computation - Where possible, the architecture aims to only process values as required to maintain correctness.

Localized memory - Each module is responsible for storing its own configuration and state using localized SRAM memories. No global or off-chip memory is used.

Architecture Overview



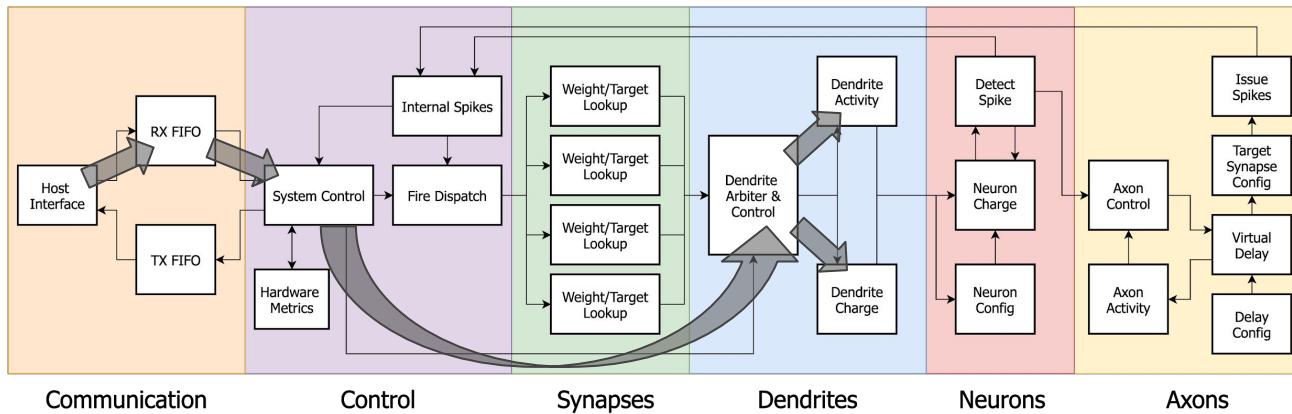
Architecture Overview



Step 1: The host sends commands to configure a network.

The arrows indicate where the network configuration is stored.

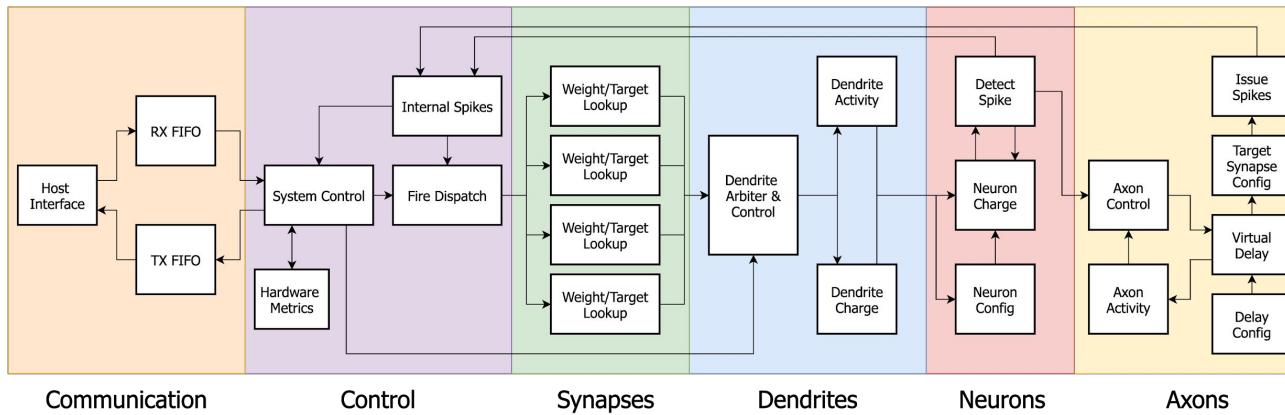
Architecture Overview



Step 2: The host sends inputs for the network which are accumulated into the dendrite units.

The arrows indicate the flow of data from host to dendrite

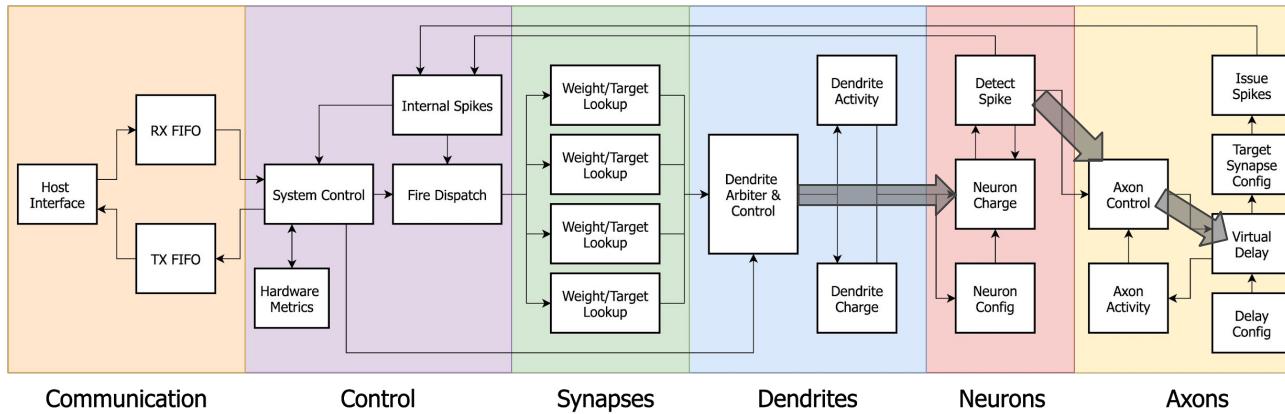
Architecture Overview



Step 3: The host sends a step forward command to advance time by t steps.

The host may send additional inputs in between step commands.

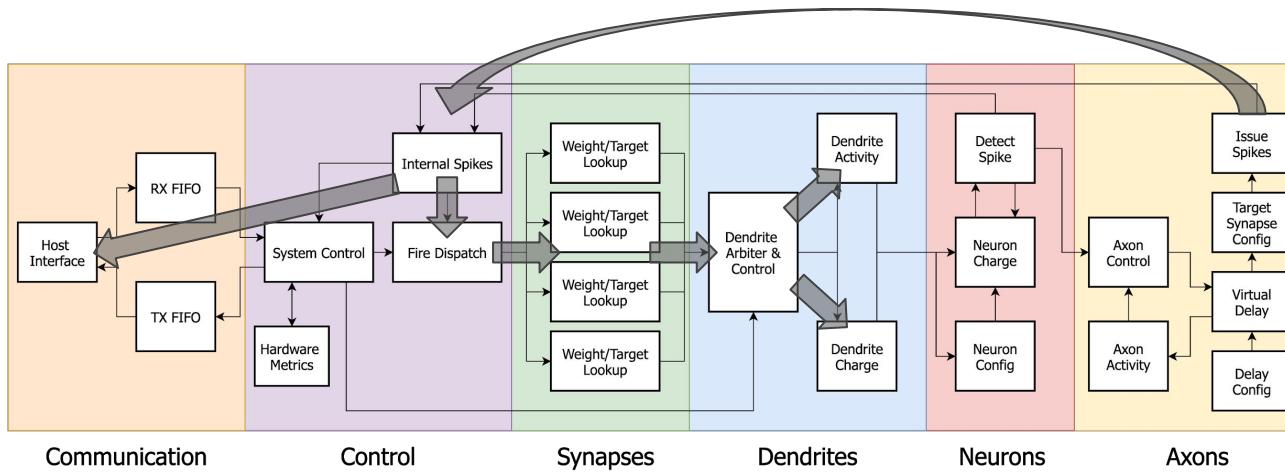
Architecture Overview



Step 4a: The first phase of network execution

Any stored charge in the dendrites is flushed to the neurons.
The neurons update state and determine if it should spike.
Neuron spikes are sent the axon unit.

Architecture Overview



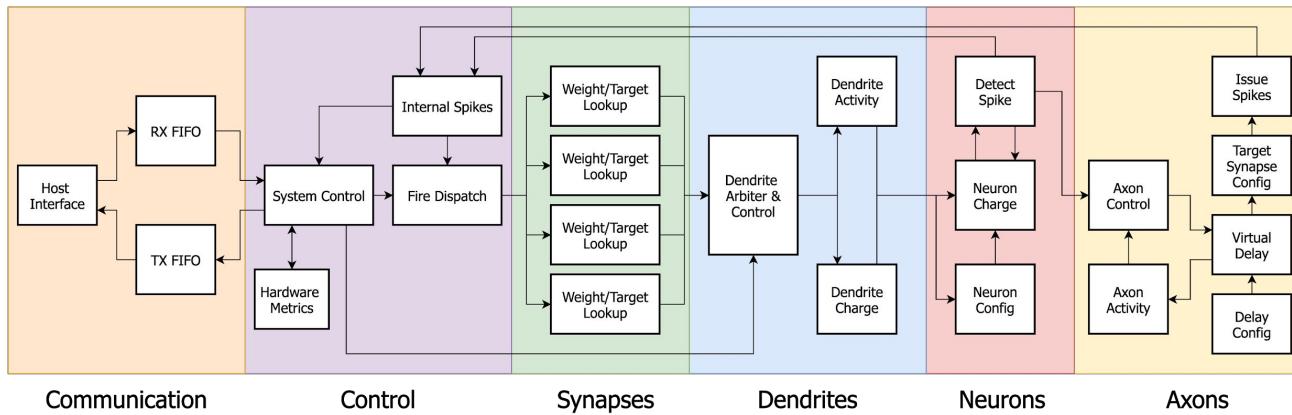
Step 4b: The second phase of network execution

The axon unit flushes any delayed spikes for the new timestep.

For each axon spike, a set of synapses will look up value/target pairs.

The dendrite unit will accumulate the value for the target neuron address.

Architecture Overview



Steps 2-4 repeat until the network has received all inputs and processed all desired time steps.

Note: Steps 4a and 4b overlap in execution during each timestep.

Architecture Overview

Steps to running a network:

- Step 1: Configure the neurons and synapses
- Step 2: Host sends input to the network
- Step 3: Advance time
- Step 4: Process all activity at time t
 - Flush charge from dendrite buffers to the neuron
 - Update neuron state and spike if charge > threshold
 - Axon accepts spikes from the neuron and tracks delay using virtual shift register in SRAM
 - If remaining delay is zero, transmit range of synapse to spike to the spike dispatch
 - Dispatch spikes to synapse units which will look up the weight & target neuron
 - Accumulate new charge from synapses into the dendrite buffers
- Repeat steps 2-4 as necessary until network is at the target time

A note about activity tracking

- Both dendrites and axons have activity tracking to reduce unnecessary work
- Each unique dendrite/axon is grouped into sets of 16
- A 16-bit register (16 bits * 16 per group) represents activity of each group
- Only groups with a set bit are evaluated for a timestep
- This approach is a balance of low area overhead and extra state processing
- Neuron activity is triggered by dendrite activity so tracking is not needed
- Synapse activity is triggered by axon activity so tracking is not needed

Communication with Host

- Dual 8-bit AXI4-Stream interfaces to the µCaspian core
- 8x512 FIFO for both send and receive channels
- Variable length packet format for minimizing required bandwidth
- Command packets are coupled to the core controller's state machine
- Output spikes & command responses are sent through a separate state machine

Variable Length Packets

Command Packets

- Add Charge (2 bytes)
- Step Time (2 bytes)
- Get Metric (2 bytes)
- Clear Activity (1 byte)
- Clear Configuration (1 byte)
- Configure Neuron (7 bytes)
- Configure Synapse (5 bytes)
- Configure Multiple Synapses (5 bytes + 2 bytes per synapse)

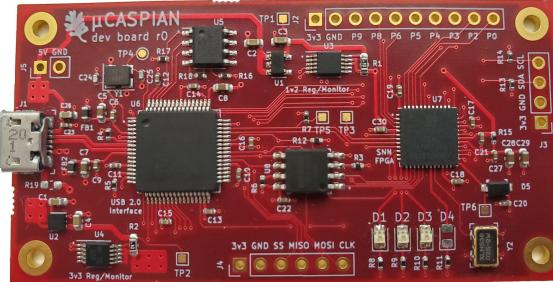
Response Packets

- Configuration Acknowledgement (1 byte)
- Clear Acknowledgement (1 byte)
- Metric Response (3 bytes)
- Time Update (5 bytes)
- Output Spike (2 bytes)

μ Caspian Development Board

The Development Board

- 2.75" by 1.4" custom PCB design for Caspian development
- Includes Lattice iCE40 UltraPlus FPGA and FT2232H USB 2.0 HS bridge
- Can be reconfigured during runtime with new FPGA bitstreams
- On-board power monitoring for 3.3V (I/O) and 1.2V (FPGA Logic)
- Exposed I/O for configuration, application data, and power monitoring
- Could be built at volume for less than \$35



Target FPGA: Lattice ice40 up5k

- 5280 Logic Cells consisting of a 4-input LUT and a DFF
- 30 4Kbit pseudo dual port SRAMs (16x256)
- 4 256Kbit single port SRAMs (16x16384)
- 8 MAC units (32 bit accumulator or 16 bit multiply-accumulate)
- Free closed source toolchain: Lattice Radiant
- Open source toolchain: yosys + nextpnr + icestorm

Target Communications: FTDI 2232H

- USB 2.0 High Speed Bridge Chip
- Two independent communication channels
- Channel A - FT245 Asynchronous FIFO for host to FPGA communication
- Channel B - SPI for configuring the bitstream on the ice40 FPGA
- Open source software library: libftdi + libusb

Annotated Development Board

